


SQL Window Functions Cheat Sheet


WINDOW FUNCTIONS

compute their result based on a sliding window frame, a set of rows that are somehow related to the current row.



AGGREGATE FUNCTIONS VS. WINDOW FUNCTIONS

unlike aggregate functions, window functions do not collapse rows.



SYNTAX

```
SELECT city, month,
       sum(sold) OVER (
         PARTITION BY city
         ORDER BY month
         RANGE UNBOUNDED PRECEDING) total
FROM sales;
```

```
SELECT <column_1>, <column_2>,
       <window_function>() OVER (
         PARTITION BY <...>
         ORDER BY <...>
         <window_frame>) <window_column_alias>
FROM <table_name>;
```

Named Window Definition

```
SELECT country, city,
       rank() OVER country_sold_avg
FROM sales
WHERE month BETWEEN 1 AND 6
GROUP BY country, city
HAVING sum(sold) > 10000
WINDOW country_sold_avg AS (
  PARTITION BY country
  ORDER BY avg(sold) DESC)
ORDER BY country, city;
```

```
SELECT <column_1>, <column_2>,
       <window_function>() OVER <window_name>
FROM <table_name>
WHERE <...>
GROUP BY <...>
HAVING <...>
WINDOW <window_name> AS (
  PARTITION BY <...>
  ORDER BY <...>
  <window_frame>)
ORDER BY <...>;
```

PARTITION BY, ORDER BY, and window frame definition are all optional.

LOGICAL ORDER OF OPERATIONS IN SQL

1. FROM, JOIN

2. WHERE

3. GROUP BY

4. aggregate functions

5. HAVING

6. window functions
7. SELECT

8. DISTINCT

9. UNION/INTERSECT/EXCEPT

10. ORDER BY

11. OFFSET

12. LIMIT/FETCH/TOP

You can use window functions in SELECT and ORDER BY. However, you can't put window functions anywhere in the FROM, WHERE, GROUP BY, or HAVING clauses.

PARTITION BY

divides rows into multiple groups, called **partitions**, to which the window function is applied.

| PARTITION BY city | | | | |
|-------------------|--------|------|-----|-----|
| month | city | sold | | sum |
| 1 | Rome | 200 | | |
| 2 | Paris | 500 | | |
| 1 | London | 100 | | |
| 1 | Paris | 300 | | |
| 2 | Rome | 300 | | |
| 2 | London | 400 | | |
| 3 | Rome | 400 | | |
| 1 | Paris | 300 | 800 | |
| 2 | Paris | 500 | 800 | |
| 1 | Rome | 200 | 900 | |
| 2 | Rome | 300 | 900 | |
| 3 | Rome | 400 | 900 | |
| 1 | London | 100 | 500 | |
| 2 | London | 400 | 500 | |

Default Partition: with no PARTITION BY clause, the entire result set is the partition.

ORDER BY

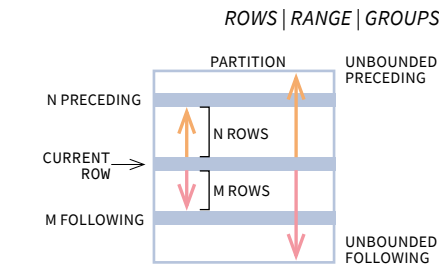
specifies the order of rows in each partition to which the window function is applied.

| PARTITION BY city ORDER BY month | | | | |
|----------------------------------|--------|-------|--|--|
| sold | city | month | | |
| 200 | Rome | 1 | | |
| 500 | Paris | 2 | | |
| 100 | London | 1 | | |
| 300 | Paris | 1 | | |
| 300 | Rome | 2 | | |
| 400 | London | 2 | | |
| 400 | Rome | 3 | | |
| 300 | Paris | 1 | | |
| 500 | Paris | 2 | | |
| 200 | Rome | 1 | | |
| 300 | Rome | 2 | | |
| 400 | Rome | 3 | | |
| 100 | London | 1 | | |
| 400 | London | 2 | | |

Default ORDER BY: with no ORDER BY clause, the order of rows within each partition is arbitrary.

WINDOW FRAME

is a set of rows that are somehow related to the current row. The window frame is evaluated separately within each partition.



The bounds can be any of the five options:

- UNBOUNDED PRECEDING
- n PRECEDING
- CURRENT ROW
- n FOLLOWING
- UNBOUNDED FOLLOWING

The lower_bound must be BEFORE the upper_bound

| ROWS BETWEEN 1 PRECEDING AND 1 FOLLOWING | | | |
|--|------|-------|--|
| city | sold | month | |
| Paris | 300 | 1 | |
| Rome | 200 | 1 | |
| Paris | 500 | 2 | |
| Rome | 100 | 4 | |
| Paris | 200 | 4 | |
| Paris | 300 | 5 | |
| Rome | 200 | 5 | |
| London | 200 | 5 | |
| London | 100 | 6 | |
| Rome | 300 | 6 | |
| 1 row before the current row and 1 row after the current row | | | |

| RANGE BETWEEN 1 PRECEDING AND 1 FOLLOWING | | | |
|---|------|-------|--|
| city | sold | month | |
| Paris | 300 | 1 | |
| Rome | 200 | 1 | |
| Paris | 500 | 2 | |
| Rome | 100 | 4 | |
| Paris | 200 | 4 | |
| Paris | 300 | 5 | |
| Rome | 200 | 5 | |
| London | 200 | 5 | |
| London | 100 | 6 | |
| Rome | 300 | 6 | |
| values in the range between 3 and 5 | | | |
| ORDER BY must contain a single expression | | | |

| GROUPS BETWEEN 1 PRECEDING AND 1 FOLLOWING | | | |
|--|------|-------|--|
| city | sold | month | |
| Paris | 300 | 1 | |
| Rome | 200 | 1 | |
| Paris | 500 | 2 | |
| Rome | 100 | 4 | |
| Paris | 200 | 4 | |
| Paris | 300 | 5 | |
| Rome | 200 | 5 | |
| London | 200 | 5 | |
| London | 100 | 6 | |
| Rome | 300 | 6 | |
| 1 group before the current row and 1 group after the current row regardless of the value | | | |

As of 2020, GROUPS is only supported in PostgreSQL 11 and up.

ABBREVIATIONS

| Abbreviation | Meaning |
|---------------------|---|
| UNBOUNDED PRECEDING | BETWEEN UNBOUNDED PRECEDING AND CURRENT ROW |
| n PRECEDING | BETWEEN n PRECEDING AND CURRENT ROW |
| CURRENT ROW | BETWEEN CURRENT ROW AND CURRENT ROW |
| n FOLLOWING | BETWEEN AND CURRENT ROW AND n FOLLOWING |
| UNBOUNDED FOLLOWING | BETWEEN CURRENT ROW AND UNBOUNDED FOLLOWING |

DEFAULT WINDOW FRAME

If ORDER BY is specified, then the frame is RANGE BETWEEN UNBOUNDED PRECEDING AND CURRENT ROW.
Without ORDER BY, the frame specification is ROWS BETWEEN UNBOUNDED PRECEDING AND UNBOUNDED FOLLOWING.

SQL Window Functions Cheat Sheet

LIST OF WINDOW FUNCTIONS

Aggregate Functions

- `avg()`
- `count()`
- `max()`
- `min()`
- `sum()`

Ranking Functions

- `row_number()`
- `rank()`
- `dense_rank()`

Distribution Functions

- `percent_rank()`
- `cume_dist()`

Analytic Functions

- `lead()`
- `lag()`
- `ntile()`
- `first_value()`
- `last_value()`
- `nth_value()`

AGGREGATE FUNCTIONS

- `avg(expr)` – average value for rows within the window frame
- `count(expr)` – count of values for rows within the window frame
- `max(expr)` – maximum value within the window frame
- `min(expr)` – minimum value within the window frame
- `sum(expr)` – sum of values within the window frame

ORDER BY and Window Frame: Aggregate functions do not require an ORDER BY. They accept window frame definition (ROWS, RANGE, GROUPS).

RANKING FUNCTIONS

- `row_number()` – unique number for each row within partition, with different numbers for tied values
- `rank()` – ranking within partition, with gaps and same ranking for tied values
- `dense_rank()` – ranking within partition, with no gaps and same ranking for tied values

| city | price | row_number | rank | dense_rank |
|--------|-------|----------------------|------|------------|
| | | over(order by price) | | |
| Paris | 7 | 1 | 1 | 1 |
| Rome | 7 | 2 | 1 | 1 |
| London | 8.5 | 3 | 3 | 2 |
| Berlin | 8.5 | 4 | 3 | 2 |
| Moscow | 9 | 5 | 5 | 3 |
| Madrid | 10 | 6 | 6 | 4 |
| Oslo | 10 | 7 | 6 | 4 |

ORDER BY and Window Frame: `rank()` and `dense_rank()` require ORDER BY, but `row_number()` does not require ORDER BY. Ranking functions do not accept window frame definition (ROWS, RANGE, GROUPS).

ANALYTIC FUNCTIONS

- `lead(expr, offset, default)` – the value for the row *offset* rows after the current; *offset* and *default* are optional; default values: *offset* = 1, *default* = NULL
- `lag(expr, offset, default)` – the value for the row *offset* rows before the current; *offset* and *default* are optional; default values: *offset* = 1, *default* = NULL

lead(sold) OVER(ORDER BY month)

| month | sold | |
|-------|------|------|
| 1 | 500 | 300 |
| 2 | 300 | 400 |
| 3 | 400 | 100 |
| 4 | 100 | 500 |
| 5 | 500 | NULL |

lag(sold) OVER(ORDER BY month)

| month | sold | |
|-------|------|------|
| 1 | 500 | NULL |
| 2 | 300 | 500 |
| 3 | 400 | 300 |
| 4 | 100 | 400 |
| 5 | 500 | 100 |

lead(sold, 2, 0) OVER(ORDER BY month)

| month | sold | |
|-------|------|-----|
| 1 | 500 | 400 |
| 2 | 300 | 100 |
| 3 | 400 | 500 |
| 4 | 100 | 0 |
| 5 | 500 | 0 |

lag(sold, 2, 0) OVER(ORDER BY month)

| month | sold | |
|-------|------|-----|
| 1 | 500 | 0 |
| 2 | 300 | 0 |
| 3 | 400 | 500 |
| 4 | 100 | 300 |
| 5 | 500 | 400 |

- `ntile(n)` – divide rows within a partition as equally as possible into *n* groups, and assign each row its group number.

ntile(3)

| city | sold | |
|--------|------|---|
| Rome | 100 | 1 |
| Paris | 100 | 1 |
| London | 200 | 1 |
| Moscow | 200 | 2 |
| Berlin | 200 | 2 |
| Madrid | 300 | 2 |
| Oslo | 300 | 3 |
| Dublin | 300 | 3 |

ORDER BY and Window Frame: `ntile()`, `lead()`, and `lag()` require an ORDER BY. They do not accept window frame definition (ROWS, RANGE, GROUPS).

DISTRIBUTION FUNCTIONS

- `percent_rank()` – the percentile ranking number of a row—a value in [0, 1] interval: $(rank - 1) / (total\ number\ of\ rows - 1)$
- `cume_dist()` – the cumulative distribution of a value within a group of values, i.e., the number of rows with values less than or equal to the current row's value divided by the total number of rows; a value in (0, 1] interval

percent_rank() OVER(ORDER BY sold)

| city | sold | percent_rank |
|--------|------|--------------|
| Paris | 100 | 0 |
| Berlin | 150 | 0.25 |
| Rome | 200 | 0.5 |
| Moscow | 200 | 0.5 |
| London | 300 | 1 |

without this row 50% of values are less than this row's value

cume_dist() OVER(ORDER BY sold)

| city | sold | cume_dist |
|--------|------|-----------|
| Paris | 100 | 0.2 |
| Berlin | 150 | 0.4 |
| Rome | 200 | 0.8 |
| Moscow | 200 | 0.8 |
| London | 300 | 1 |

80% of values are less than or equal to this one

ORDER BY and Window Frame: Distribution functions require ORDER BY. They do not accept window frame definition (ROWS, RANGE, GROUPS).

- `first_value(expr)` – the value for the first row within the window frame
- `last_value(expr)` – the value for the last row within the window frame

first_value(sold) OVER
(PARTITION BY city ORDER BY month)

| city | month | sold | first_value |
|-------|-------|------|-------------|
| Paris | 1 | 500 | 500 |
| Paris | 2 | 300 | 500 |
| Paris | 3 | 400 | 500 |
| Rome | 2 | 200 | 200 |
| Rome | 3 | 300 | 200 |
| Rome | 4 | 500 | 200 |

last_value(sold) OVER
(PARTITION BY city ORDER BY month
RANGE BETWEEN UNBOUNDED PRECEDING
AND UNBOUNDED FOLLOWING)

| city | month | sold | last_value |
|-------|-------|------|------------|
| Paris | 1 | 500 | 400 |
| Paris | 2 | 300 | 400 |
| Paris | 3 | 400 | 400 |
| Rome | 2 | 200 | 500 |
| Rome | 3 | 300 | 500 |
| Rome | 4 | 500 | 500 |

Note: You usually want to use RANGE BETWEEN UNBOUNDED PRECEDING AND UNBOUNDED FOLLOWING with `last_value()`. With the default window frame for ORDER BY, RANGE UNBOUNDED PRECEDING, `last_value()` returns the value for the current row.

- `nth_value(expr, n)` – the value for the *n*-th row within the window frame; *n* must be an integer

nth_value(sold, 2) OVER (PARTITION BY city
ORDER BY month RANGE BETWEEN UNBOUNDED
PRECEDING AND UNBOUNDED FOLLOWING)

| city | month | sold | nth_value |
|--------|-------|------|-----------|
| Paris | 1 | 500 | 300 |
| Paris | 2 | 300 | 300 |
| Paris | 3 | 400 | 300 |
| Rome | 2 | 200 | 300 |
| Rome | 3 | 300 | 300 |
| Rome | 4 | 500 | 300 |
| Rome | 5 | 300 | 300 |
| London | 1 | 100 | NULL |

ORDER BY and Window Frame: `first_value()`, `last_value()`, and `nth_value()` do not require an ORDER BY. They accept window frame definition (ROWS, RANGE, GROUPS).